

---

# IDS 435 - Optimization via Gurobi (Part 2)

Spring 2022

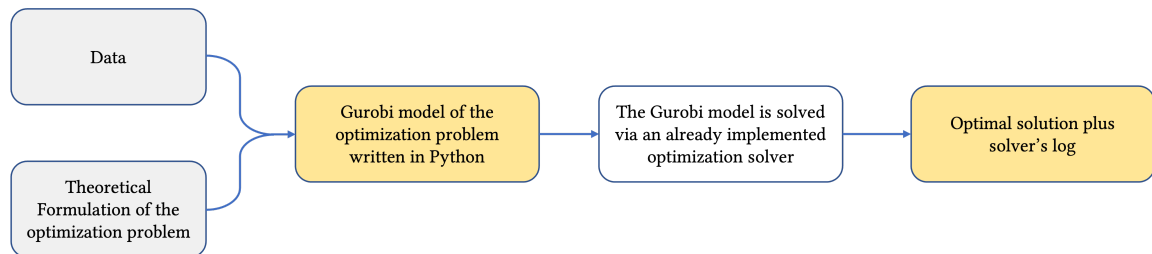
---

## Table of Contents

- [A Procedure to Model an Optimization Problem in Gurobi](#)
  - [A Toy Example](#)
  - [Simple Linear Regression](#)
  - [Feature Selection in Regression](#)
- 

## A Procedure to Model an Optimization Problem in Gurobi

The main task to use Gurobi in practice is to convert the mathematical formulation of an optimization problem into a Gurobi model that can be solved efficiently.



### The Procedure

#### 1. Sets and indices

Use python `list`, `range`, `arrange`, etc to define index sets to count over decision variables and constraints.

#### 2. Parameters (i.e., data)

Data should be formatted in a way that is easy to define decision variables, objective function, and constraints in Gurobi. Data can be represented in different formats such as

- Python lists
- Python dictionaries
- Numpy array

- Pandas `dataframe`
- Gurobi multidict

Since Numpy arrays are widely-used data structures in python, we employ them to represent data that later on used in a Gurobi model

### 3. Decision variables

- Specify type of the variable (real-valued, binary, integer, etc)
- Specify if the variable is signed or not (positive or negative)
- Try to define put related decision variables of the model in an array of variables instead of defining them individually

```
# For loop
for i in range (N):
    model.addVar( ... )
```

```
# One line definition
model.addMvar( shape=N)
```

### 4. Constraints

- Specify type of a constraint (linear, quadratic, etc)
- Find an appropriate Gurobi function to model the constraint (oftentimes you can use `addConstrs` )
- Try to define multiple constraints in a single line of code via python inline for loop

```
# For loop
for i in range (N):
    model.addConstr( ... )
```

```
# Inline for loop (preferred)
model.addConstrs( ... for i in range (N))
```

### 5. Objective function

- Define the objective function using data and decision variables
- Specify if the problem is maximization or minimization

### 6. Optimize

- Choose a solver
- Specify parameters of the solver (i.e., stopping criteria, feasibility tolerance, etc)
- Solve the Gurobi model

### 7. Analyze results (*Gurobi solved the model*)

- Is model well-conditioned (i.e., no numerical issues encountered while optimization)?



```

"""          Step 4. Constraints          """
model.addConstrs(gb.quicksum(constr_matrix[i][j]*x[j] for j in var_index[i]))

"""          Step 5. Objective function          """
model.setObjective(-gb.quicksum(x))

"""          Step 6. Optimize          """
print('='*100)
model.setParam('Method',2)
model.setParam('Crossover',0)
model.update()
model.optimize()
print('='*100)

"""          Step 7. Analyze results          """
print('The optimal x_1:      \t',      x[0].X)
print('The optimal x_2:      \t',      x[1].X)
print('The optimal value is:  \t',      model.ObjVal)
print('='*100)

```

```

Set parameter Username
Academic license - for non-commercial use only - expires 2022-05-07
=====
=====
Set parameter Method to value 2
Set parameter Crossover to value 0
Gurobi Optimizer version 9.5.1 build v9.5.1rc2 (mac64[rosetta2])
Thread count: 10 physical cores, 10 logical processors, using up to 10 threads
Optimize a model with 2 rows, 2 columns and 4 nonzeros
Model fingerprint: 0xed33d8fe
Coefficient statistics:
  Matrix range      [1e+00, 2e+00]
  Objective range   [1e+00, 1e+00]
  Bounds range      [0e+00, 0e+00]
  RHS range         [1e+00, 1e+00]
Presolve time: 0.00s
Presolved: 2 rows, 2 columns, 4 nonzeros
Ordering time: 0.00s

Barrier statistics:
AA' NZ      : 1.000e+00
Factor NZ   : 3.000e+00
Factor Ops  : 5.000e+00 (less than 1 second per iteration)
Threads     : 1

      Objective          Residual
Iter   Primal      Dual      Primal    Dual      Compl      Time
  0  -8.67927042e-01 -4.61538462e-01  1.51e-01 3.08e-01  2.86e-01    0s
  1  -6.05231787e-01 -6.96010401e-01  0.00e+00 0.00e+00  2.27e-02    0s
  2  -6.66536989e-01 -6.66799107e-01  0.00e+00 0.00e+00  6.55e-05    0s
  3  -6.66666537e-01 -6.66666799e-01  0.00e+00 0.00e+00  6.55e-08    0s
  4  -6.66666667e-01 -6.66666667e-01  0.00e+00 2.22e-16  6.55e-11    0s

Barrier solved model in 4 iterations and 0.01 seconds (0.00 work units)
Optimal objective -6.66666667e-01

=====
=====
The optimal x_1:          0.3333333332685205
The optimal x_2:          0.3333333332685205
The optimal value is:     -0.666666666537041
=====
=====

```

---

## Simple Linear Regression

This application is motivated by [curve fitting](#) Gurobi example.

Consider a dataset of  $N$  points  $\{(x^i, y^i) : i = 1, 2, \dots, N\}$ . We want to fit a linear model with intercept  $\beta_0$  and slope  $\beta_1$  to approximate response variable  $y^i$  using feature  $x^i$ . Specifically, we want the following constraints to hold as close as possible:

$$y^i \approx \beta_0 + \beta_1 x^i, \quad \forall i = 1, 2, \dots, N.$$

For each  $i$ ,  $\beta_0 + \beta_1 x^i$  is the approximate value for  $y_i$ . Our objective is to compute  $\beta_0$  and  $\beta$  such that an "error" is minimized. Below we discuss two metrics of error and implement them in Gurobi.

1. Mean Absolute Deviation (MAD; *linear objective function*):

$$\begin{aligned} \min_{\beta_0, \beta_1} \quad & \frac{1}{N} \sum_{i=1}^N (u_i + v_i) \\ & y^i = \beta_0 + \beta_1 x^i + u_i - v_i, \quad \forall i = 1, 2, \dots, N, \\ & u_i, v_i \geq 0, \quad \forall i = 1, 2, \dots, N, \\ & \beta_0, \beta_1 \text{ unrestricted.} \end{aligned}$$

2. Mean Squared Error (MSE; *quadratic objective function*):

$$\begin{aligned} \min_{\beta_0, \beta_1} \quad & \frac{1}{N} \sum_{i=1}^N \varepsilon_i^2 \\ & \varepsilon_i = y^i - \beta_0 - \beta_1 x^i, \quad \forall i = 1, 2, \dots, N, \\ & \varepsilon_i \text{ unrestricted,} \quad \forall i = 1, 2, \dots, N, \\ & \beta_0, \beta_1 \text{ unrestricted.} \end{aligned}$$

Gurobi model for MAD minimization

```
In [ ]: import numpy as np
import gurobi as gb
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

if __name__ == "__main__":

    """      Step 1. Sets and indices      """
    num_data      = 10
    data_index     = range(num_data)

    """      Step 2. Parameters      """
    x_train        = np.array([0., 0.5, 1, 1.5, 1.9, 2.5,
                               4.5, 5., 5.5, 6, 6.6, 7,
    y_train         = np.array([1., 0.9, .7, 1.5, 2., 2.4,
                               3.5, 1., 4, 3.6, 2.7, 5.

    """      Step 3. Decision Variables      """
    model          = gb.Model('MAD')

    ## coef := (\beta_0, \beta)
    coef           = model.addMVar(
```

```

        shape      = 2,
        vtype      = gb.GRB.CONTINUOUS,
        lb         = -gb.GRB.INFINITY,
        ub         =  gb.GRB.INFINITY )

u          = model.addMVar(
        shape      = num_data,
        vtype      = gb.GRB.CONTINUOUS,
        lb         = 0,
        ub         = gb.GRB.INFINITY )

v          = model.addMVar(
        shape      = num_data,
        vtype      = gb.GRB.CONTINUOUS,
        lb         = 0,
        ub         = gb.GRB.INFINITY )

"""          Step 4. Constraints          """
model.addConstrs(y_train[i] == coef[0] + coef[1]*x_train[i] + u[i] - v[i])

"""          Step 5. Objective function          """
model.setObjective(gb.quicksum(u) + gb.quicksum(v))

"""          Step 6. Optimize          """
print('='*100)
model.update()
model.optimize()
print('='*100)

"""          Step 7. Analyze results          """
MAD_intercept, MAD_slope = coef[0].X,coef[1].X
print('The optimal intercept:\t', MAD_intercept)
print('The optimal slope:      \t', MAD_slope)
print('The optimal objective value: \t', model.ObjVal)
print('='*100)

fig = plt.figure(figsize=(8,6))
plt.plot(x_train,y_train,marker='o',lw=1,color='black',label='True Data')
plt.plot(x_train,[MAD_intercept + x*MAD_slope for x in x_train],color='red',lw=2)
plt.legend(fontsize=12)
plt.show()

print([np.round(u[i].X,4) for i in data_index])
print([np.round(v[i].X,4) for i in data_index])

```

```

=====
=====
Gurobi Optimizer version 9.5.1 build v9.5.1rc2 (mac64[rosetta2])
Thread count: 10 physical cores, 10 logical processors, using up to 10 threads
Optimize a model with 10 rows, 22 columns and 39 nonzeros
Model fingerprint: 0x1a0d94ca
Coefficient statistics:
  Matrix range      [5e-01, 4e+00]
  Objective range   [1e+00, 1e+00]
  Bounds range      [0e+00, 0e+00]
  RHS range         [7e-01, 4e+00]
Presolve time: 0.00s
Presolved: 10 rows, 22 columns, 39 nonzeros

```

Iteration	Objective	Primal Inf.	Dual Inf.	Time
0	handle free variables			0s
11	3.3900000e+00	0.000000e+00	0.000000e+00	0s

```

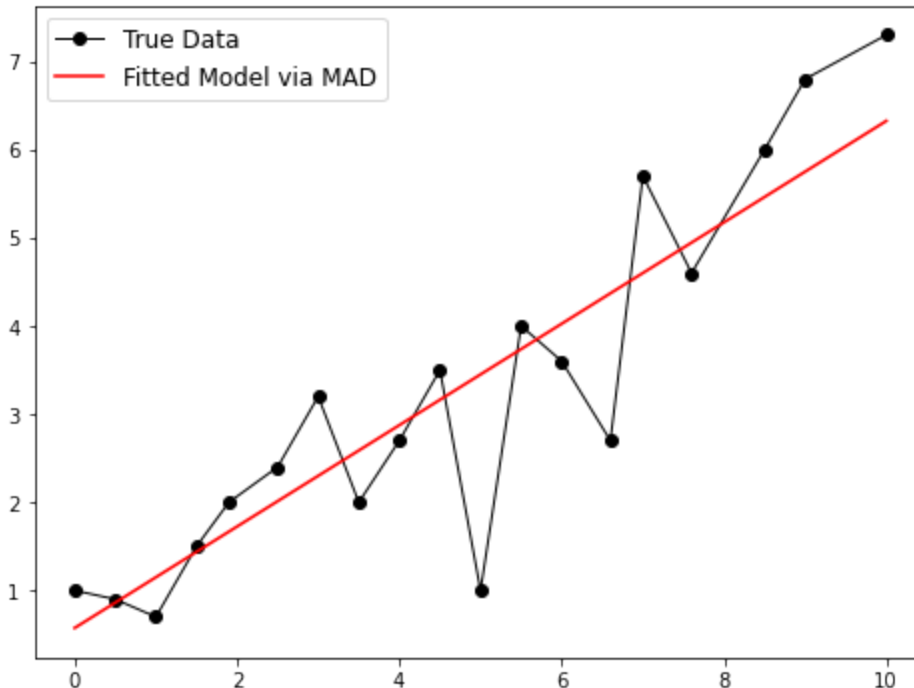
Solved in 11 iterations and 0.01 seconds (0.00 work units)
Optimal objective 3.390000000e+00
=====
=====

```

```

The optimal intercept: 0.575
The optimal slope:    0.65
The optimal objective value: 3.3899999999999997
=====
=====

```



```

[0.425, 0.0, 0.0, 0.0, 0.19, 0.2, 0.675, 0.0, 0.0, 0.0]
[0.0, 0.0, 0.525, 0.05, 0.0, 0.0, 0.0, 0.85, 0.475, 0.0]

```

Gurobi model for MSE minimization

```

In [ ]: import numpy as np
import gurobi as gb

```



```

import matplotlib.pyplot as plt

if __name__ == "__main__":

    """        Step 1. Sets and indices        """
    num_data      = 10
    data_index    = range(num_data)

    """        Step 2. Parameters        """
    x_train       = np.array([0., 0.5, 1, 1.5, 1.9, 2.5,
                              4.5, 5., 5.5, 6, 6.6, 7,
                              ])
    y_train       = np.array([1., 0.9, .7, 1.5, 2., 2.4,
                              3.5, 1., 4, 3.6, 2.7, 5,
                              ])

    """        Step 3. Decision variables        """

    model         = gb.Model('MSE')

    ## coef := (\beta_0, \beta)
    coef          = model.addMVar(
        shape      = 2,
        vtype      = gb.GRB.CONTINUOUS,
        lb         = -gb.GRB.INFINITY,
        ub         = gb.GRB.INFINITY )

    diff          = model.addMVar(
        shape      = num_data,
        vtype      = gb.GRB.CONTINUOUS,
        lb         = -gb.GRB.INFINITY,
        ub         = gb.GRB.INFINITY )

    """        Step 4. Constraints        """
    model.addConstr(diff[i] == y_train[i] - coef[0] - coef[1]*x_train[i]

    """        Step 5. Objective function        """
    model.setObjective(gb.quicksum(diff[i]*diff[i] for i in data_index))

    """        Step 6. Optimize        """
    print('='*100)
    model.update()
    model.optimize()
    print('='*100)

    """        Step 7. Analyze results        """
    MSE_intercept, MSE_slope = coef[0].X, coef[1].X
    print('The optimal intercept:\t', MSE_intercept)
    print('The optimal slope: \t', MSE_slope)
    print('The optimal objective value: \t', model.ObjVal)
    print('='*100)

    fig = plt.figure(figsize=(8,6))
    plt.plot(x_train,y_train,marker='o',lw=1,color='black',label='True Data')
    plt.plot(x_train,[MAD_intercept + x*MAD_intercept for x in x_train],color='red',label='MAD Model')

```

```
plt.plot(x_train, [MSE_intercept + x*MSE_slope for x in x_train], color='b')
plt.legend(fontsize=12)
plt.show()
```

```
=====
=====
Gurobi Optimizer version 9.5.1 build v9.5.1rc2 (mac64[rosetta2])
Thread count: 10 physical cores, 10 logical processors, using up to 10 threads
```

```
Optimize a model with 10 rows, 12 columns and 29 nonzeros
```

```
Model fingerprint: 0x2ce89fab
```

```
Model has 10 quadratic objective terms
```

```
Coefficient statistics:
```

```
Matrix range      [5e-01, 4e+00]
```

```
Objective range   [0e+00, 0e+00]
```

```
QObjective range  [2e+00, 2e+00]
```

```
Bounds range      [0e+00, 0e+00]
```

```
RHS range         [7e-01, 4e+00]
```

```
Presolve removed 1 rows and 1 columns
```

```
Presolve time: 0.00s
```

```
Presolved: 9 rows, 11 columns, 27 nonzeros
```

```
Presolved model has 10 quadratic objective terms
```

```
Ordering time: 0.00s
```

```
Barrier statistics:
```

```
Free vars   : 11
```

```
AA' NZ      : 3.600e+01
```

```
Factor NZ   : 4.500e+01
```

```
Factor Ops  : 2.850e+02 (less than 1 second per iteration)
```

```
Threads     : 1
```

Iter	Objective		Residual		Compl	Time
	Primal	Dual	Primal	Dual		
0	1.78781526e+00	-1.78781526e+00	3.77e-15	7.83e-01	0.00e+00	0s
1	1.78685941e+00	1.78685583e+00	1.13e-08	7.83e-07	0.00e+00	0s
2	1.78685941e+00	1.78685941e+00	2.30e-14	7.82e-13	0.00e+00	0s

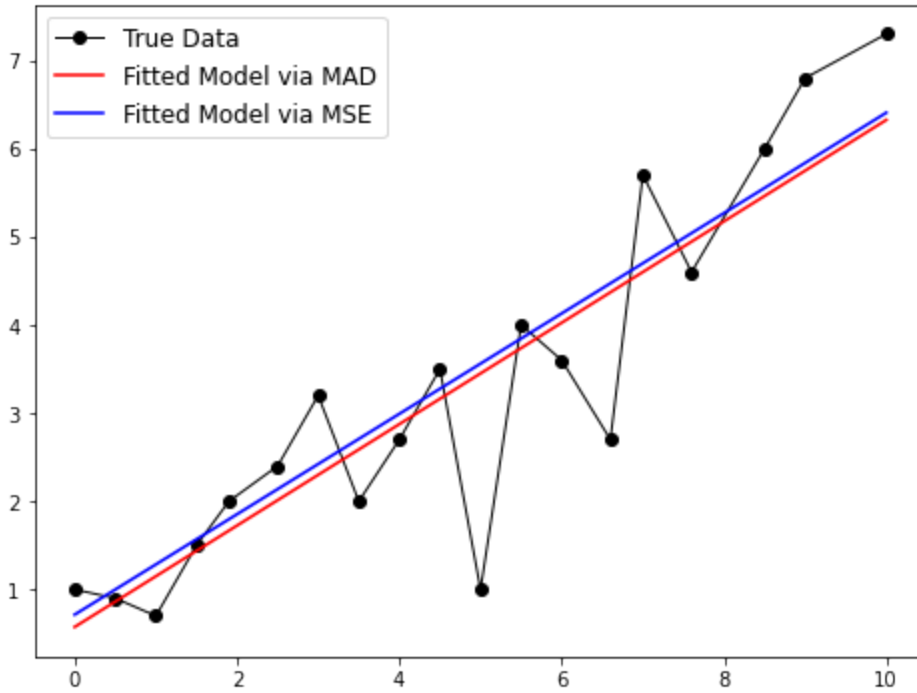
```
Barrier solved model in 2 iterations and 0.01 seconds (0.00 work units)
```

```
Optimal objective 1.78685941e+00
```

```
=====
=====
The optimal intercept: 0.7149197447302484
```

```
The optimal slope: 0.5692322568168443
```

```
The optimal objective value: 1.7868594082382703
```



## Feature Selection in Regression

This application is motivated by [Feature Selection for Forecasting](#) Gurobi example.

### Preliminaries

Building on the previous example, we study a linear regression problem in which the optimized linear model should only use a small subset of features. In other words, we need to select the best subset of features that their linear combinations have the lowest training error.

Consider training set  $\{(x^i, y^i), i = 1, 2, \dots, N\}$  with  $N$  observations, where  $x^i = (x_1^i, x_2^i, \dots, x_d^i)$  and  $y^i$  are  $d$ -dimensional feature vector and the value of the response variable associated with the  $i$ -th observation, respectively.

The goal in linear regression problem is to find coefficients  $\beta = (\beta_0, \beta_1, \beta_2, \dots, \beta_d)$ , where  $\beta_0$  is the intercept and  $(\beta_1, \beta_2, \dots, \beta_d)$  is the slope of the linear model, such that the following approximation errors attain their lowest values:

$$\varepsilon_i := y^i - \left( \beta_0 + \sum_{j=1}^d \beta_j x_j^i \right), \quad \forall i = 1, 2, \dots, N.$$

**Ordinary least squares (OLS).** OLS is an optimization problem that allows us to obtain the lowest approximation error discussed above. OLS can be written as follows:

$$\min_{\beta} \text{MSE}(\beta) \tag{OLS}$$

where

$$\text{MSE}(\beta) = \frac{1}{N} \sum_{i=1}^N \varepsilon_i^2 = \frac{1}{N} \sum_{i=1}^N \left[ y^i - \left( \beta_0 + \sum_{j=1}^d \beta_j x_j^i \right) \right]^2$$

**Indirect subset selection via LASSO regression.** OLS does not perform the subset selection, that is, allowing only  $K$  out of  $d + 1$  elements of  $\beta$  take a non-zero value. The LASSO regression is a modification of OLS to indirectly perform subset selection. It is given by the following optimization problem:

$$\min_{\beta} \text{MSE}(\beta) \quad \text{s.t.} \quad \sum_{j=0}^d |\beta_j| \leq T. \quad (\text{LASSO})$$

LASSO is a convex optimization problem. LASSO does not guarantee that its optimal  $\beta^*$  has at most  $K$  non-zero elements. That is why we refer to it as an indirect approach for subset selection.

**Direct subset selection via  $l_0$ -regression.** To directly perform subset selection, we cannot stay in the convex optimization world since counting the number of non-zeros of a vector is not a convex function. The non-convex optimization problem that allows us to perform subset selection is given by the following  $L_0$ -regression problem:

$$\min_{\beta} \text{MSE}(\beta) \quad \text{s.t.} \quad \|\beta\|_0 \leq K \quad (l_0\text{-regression})$$

where

$$\|\beta\|_0 := \text{number of non-zero elements of } \beta.$$

### Comparison of OLS, LASSO, and $l_0$ -regression

	OLS	LASSO	$l_0$ -regression
Convex	Yes	Yes	No
Indirect subset selection	No	Yes	Yes
Direct subset selection	No	No	Yes
Constrained?	No	Yes	Yes
How to solve?	<code>LinearRegression</code> in sklearn	<code>Lasso</code> in sklearn	A model in Gurobi

### Reformulating $l_0$ -regression

Let's focus on formulating  $l_0$  constraint  $\|\beta\|_0 \leq K$ . This constraint can be modeled using the following constraints that are written via binary variables:

$$z_j := \begin{cases} 1 & \text{if } \beta_j \neq 0 \\ 0 & \text{if } \beta_j = 0 \end{cases}, \quad \forall j = 0, 1, \dots, d;$$

$$\sum_{j=0}^d z_j = K$$

Overall, we can reformulate  $l_0$ -regression problem as the following optimization problem with both continuous and integer variables:

$$\min_{\beta, z} \frac{1}{N} \sum_{i=1}^N \left[ y^i - \left( \beta_0 + \sum_{j=1}^d \beta_j x_j^i \right) \right]^2$$

$$z_j = 0, \quad \text{if } \beta_j = 0, \quad \forall j = 0, 1, \dots, d,$$

$$z_j = 1, \quad \text{if } \beta_j \neq 0, \quad \forall j = 0, 1, \dots, d,$$

$$\sum_{j=0}^d z_j = K,$$

$$\beta_j \text{ unrestricted}, \quad \forall j = 0, 1, \dots, d,$$

$$z_j \text{ binary}, \quad \forall j = 0, 1, \dots, d.$$

```
In [ ]: import numpy as np
import gurobi as gb

def MIQP_version_1(train_X:np.ndarray, train_y:np.ndarray, non_zero_budget:K):
    """ Step 1. Sets and indices """
    num_data,dim      = np.shape(train_X)
    data_index        = range(num_data)      # 1,2, ..., N
    dim_index          = range(dim+1)         # 0,1, ..., d+1

    """ Step 2. Parameters """
    # Data comes from the input. Append a column of ones to the
    # feature matrix train_X to account for the intercept
    train_X            = np.concatenate([np.ones((num_data, 1)),train_X], axis=0)

    """ Step 3. Decision variables """
    model              = gb.Model('MIQP')
    coef               = model.addMVar(
                        shape      = dim + 1,
                        vtype      = gb.GRB.CONTINUOUS,
                        lb         = -gb.GRB.INFINITY,
                        ub         = gb.GRB.INFINITY )

    non_zero           = model.addMVar(
                        shape      = dim+1,
                        vtype      = gb.GRB.BINARY)

    """ Step 4. Constraints """
```

```

for j in dim_index:
    model.addConstr((non_zero[j] == 0) >> (coef[j] == 0))

model.addConstr(gb.quicksum(non_zero)<=non_zero_budget)

"""        Step 5. Objective function        """
epsilon = [None for _ in data_index]
for i in data_index:
    epsilon[i] = train_y[i] - coef@train_X[i,:]

model.setObjective((1/num_data)*gb.quicksum(epsilon[i]*epsilon[i] for i

"""        Step 6. Optimize        """
model.setParam('Seed', 123)
model.update()
model.optimize()

"""        Step 7. Analyze results        """
# we let the user to analyze the results and only return the optimal sol
opt_coef = np.array([coef[j].X for j in dim_index])
opt_val = model.objVal
return opt_coef, opt_val

```

## Loading Data

```

In [ ]: from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

if __name__ == "__main__":

    housing = fetch_california_housing()

    print(housing['DESCR'])

    X = housing['data']
    y = housing['target']
    train_X, test_X, train_y, test_y = train_test_split(X, y, test_size=0.20

    # Standardize the features so they have an avg of 0 and a sample var of
    scaler = StandardScaler()
    scaler.fit(train_X)
    train_X_std = scaler.transform(train_X)
    test_X_std = scaler.transform(test_X)

```

```
.. _california_housing_dataset:
```

California Housing dataset

```
**Data Set Characteristics:**
```

```
:Number of Instances: 20640
```

```
:Number of Attributes: 8 numeric, predictive attributes and the target
```

```
:Attribute Information:
```

- MedInc median income in block group
- HouseAge median house age in block group
- AveRooms average number of rooms per household
- AveBedrms average number of bedrooms per household
- Population block group population
- AveOccup average number of household members
- Latitude block group latitude
- Longitude block group longitude

```
:Missing Attribute Values: None
```

This dataset was obtained from the StatLib repository.

[https://www.dcc.fc.up.pt/~ltorgo/Regression/cal\\_housing.html](https://www.dcc.fc.up.pt/~ltorgo/Regression/cal_housing.html)

The target variable is the median house value for California districts, expressed in hundreds of thousands of dollars (\$100,000).

This dataset was derived from the 1990 U.S. census, using one row per census block group. A block group is the smallest geographical unit for which the U.S.

Census Bureau publishes sample data (a block group typically has a population of 600 to 3,000 people).

An household is a group of people residing within a home. Since the average number of rooms and bedrooms in this dataset are provided per household, the columns may take surprisingly large values for block groups with few households and many empty houses, such as vacation resorts.

It can be downloaded/loaded using the

```
:func:`sklearn.datasets.fetch_california_housing` function.
```

```
.. topic:: References
```

- Pace, R. Kelley and Ronald Barry, Sparse Spatial Autoregressions, Statistics and Probability Letters, 33 (1997) 291-297

Trying function MIQP\_version\_1

```
In [ ]: if __name__ == "__main__":
```

```
opt_coef, opt_val = MIQP(train_X_std, train_y, non_zero_budget=2)
```

Why do we get error? Aha, we need to do step 8, the Troubleshooting.

Let's rewrite the MSE objective in Gurobi differently to make it work

Recall the definition of error  $\varepsilon_i$  and OLS that are given by

$$\varepsilon_i := y^i - \left( \beta_0 + \sum_{j=1}^d \beta_j x_j^i \right), \quad \forall i = 1, 2, \dots, N.$$

We can compactly write the above identities as follows:

$$\varepsilon := Y - \hat{X}\beta$$

where

$$\varepsilon = \begin{bmatrix} \varepsilon_1 \\ \varepsilon_2 \\ \varepsilon_3 \\ \vdots \\ \varepsilon_N \end{bmatrix}_{N \times 1}, \quad Y := \begin{bmatrix} y^1 \\ y^2 \\ y^3 \\ \vdots \\ y^N \end{bmatrix}, \quad \hat{X} := \begin{bmatrix} 1, & x_1^1, & x_2^1, & x_3^1, & \dots, & x_d^1 \\ 1, & x_1^2, & x_2^2, & x_3^2, & \dots, & x_d^2 \\ 1, & x_1^3, & x_2^3, & x_3^3, & \dots, & x_d^3 \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ 1, & x_1^N, & x_2^N, & x_3^N, & \dots, & x_d^N \end{bmatrix}_{N \times (d+1)}$$

Using the above matrix form, we can rewrite MSE as the following quadratic form that is a predefined format in Gurobi:

$$\begin{aligned} \text{MSE}(\beta) &= \frac{1}{N} \sum_{i=1}^N \varepsilon_i^2 \\ &= \frac{1}{N} (\varepsilon^\top \varepsilon) \\ &= \frac{1}{N} ((Y - \hat{X}\beta)^\top (Y - \hat{X}\beta)) \\ &= \frac{1}{N} (\beta^\top \hat{X}^\top \hat{X} \beta - 2Y^\top \hat{X} \beta + Y^\top Y) \\ &= \beta^\top \left( \frac{1}{N} \hat{X}^\top \hat{X} \right) \beta - \left( \frac{2}{N} Y^\top \hat{X} \right) \beta + \frac{1}{N} Y^\top Y \quad (\text{quadratic in } \beta) \end{aligned}$$

Overall, we get the following MIQP:



$$\min_{\beta, z} \beta^\top \left( \frac{1}{N} \hat{X}^\top \hat{X} \right) \beta - \left( \frac{2}{N} Y^\top X \right) \beta + \frac{1}{N} Y^\top Y$$

$$z_j = 0, \quad \text{if } \beta_j = 0, \quad \forall j = 0, 1, \dots, d,$$

$$z_j = 1, \quad \text{if } \beta_j \neq 0, \quad \forall j = 0, 1, \dots, d,$$

$$\sum_{j=0}^d z_j = K,$$

$$\beta_j \text{ unrestricted}, \quad \forall j = 0, 1, \dots, d.$$

$$z_j \text{ binary}, \quad \forall j = 0, 1, \dots, d.$$

```
In [ ]: import numpy as np
import gurobi as gb

def MIQP(train_X:np.ndarray, train_y:np.ndarray, non_zero_budget:int,output_
"""      Step 1. Sets and indices      """
num_data,dim      = np.shape(train_X)
data_index        = range(num_data)      # 1,2, ..., N
dim_index         = range(dim+1)         # 0,1, ..., d+1

"""      Step 2. Parameters      """
# Data comes from the input. Append a column of ones to the
# feature matrix train_X to account for the intercept
train_X           = np.concatenate([np.ones((num_data, 1)),train_X], axis=

"""      Step 3. Decision variables      """
model              = gb.Model('MIQP')
coef               = model.addMVar(
                    shape      = dim + 1,
                    vtype      = gb.GRB.CONTINUOUS,
                    lb         = -gb.GRB.INFINITY,
                    ub         = gb.GRB.INFINITY )

non_zero           = model.addMVar(
                    shape      = dim+1,
                    vtype      = gb.GRB.BINARY)

"""      Step 4. Constraints      """

for j in dim_index:
    model.addConstr((non_zero[j] == 0) >> (coef[j] == 0))

model.addConstr(gb.quicksum(non_zero)<=non_zero_budget)

"""      Step 5. Objective function      """
objective = ((coef@(train_X.T@train_X/num_data))@coef
             - (2*((train_y.T@train_X)/num_data))@coef
             + (train_y.T@train_y)/num_data)
model.setObjective(objective)

"""      Step 6. Optimize      """
```

```

Troubleshootedmodel.setParam('OutputFlag', output_flag)
model.setParam('Seed', 123)
model.update()
model.optimize()

"""      Step 7. Analyze results      """
# we let the user to analyze the results and only return the optimal sol
opt_coef = np.array([coef[j].X for j in dim_index])
opt_val = model.objVal
return opt_coef, opt_val

```

## Solving MIQP

```

In [ ]: if __name__ == "__main__":

    ## MILP using only K features
    print('='*100)
    opt_coef, opt_val = MIQP(train_X_std, train_y, non_zero_budget=5)
    print('='*100)
    print('\nMIQP Optimal Solution: \n', opt_coef)

```

## Numerical Comparison of OLS, LASSO, and MIQP

```

In [ ]: from sklearn.linear_model import LinearRegression, LassoCV
from sklearn.metrics import mean_squared_error as mse

if __name__ == "__main__":
    ## OLS regression using all features
    ols = LinearRegression()
    ols.fit(train_X_std, train_y)

    ## LASSO regression using all features
    lasso = LassoCV(cv=5)
    lasso.fit(train_X_std, train_y)

    ## MILP using only K features
    opt_coef, opt_val = MIQP(train_X_std, train_y, non_zero_budget=5, output_
test_X_std_1 = np.concatenate([np.ones((np.shape(test_X_std)[0], 1)

    print('OLS Testing MSE   :', np.round(mse(test_y, ols.predict(test_X_std_1), 2))
    print('LASSO Testing MSE :', np.round(mse(test_y, lasso.predict(test_X_std_1), 2))
    print('MIQP Testing MSE  :', np.round(mse(test_y, test_X_std_1@opt_coef), 2))

    print('\nOLS Optimal Solution: \n', np.round(ols.coef_, 2))
    print('\nLASSO Optimal Solution:\n', np.round(lasso.coef_, 2))
    print('\nMIQP Optimal Solution: \n', np.round(opt_coef, 2))

```

```

In [ ]:

```